

## Big Data Management - CS585 - Project 2

- Azharuddin Priyotomo
  - Congyuan Tang
- 

We are working under the virtual box (Ubuntu OS) and hadoop environment that the Professor provided us.

### 1. Query 1

To perform the query 1, we only need the data set from transaction file, since all output (Customer ID, Number of Transactions and Total Sum) are all available in transaction data.

First we need to load the transaction data into one variable and then group it by customer ID. From the grouped data, we select aggregate count of the transaction and the sum of transaction, and then store it to an output file in HDFS.

Query:

```
SET DEFAULT_PARALLEL 20;
trans = LOAD '/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/transactioninput'
USING PigStorage(',') as (transid:int,custid:int,
trxamt:float,noitem:int,desc:chararray);
A = group trans by custid;
B = foreach A generate group, COUNT(trans) as NumTransactions,
SUM(trans.trxamt) as TotalSum;
STORE B INTO
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/project2/query1output' USING
PigStorage();
```

Sample Output:

1	114	58570.840938568115
2	109	58381.64875411987
3	80	40175.597732543945
4	99	48407.24899673462
5	97	48511.70783042908
6	100	50382.99196243286
7	91	44795.514194488525
8	94	49123.03216457367
9	107	53491.60666656494
10	101	48595.103954315186
11	121	64296.21905517578
12	104	48114.33840751648
13	116	57267.39225959778
14	111	51052.32263946533
15	94	52909.92872238159

Description: (Customer ID, Number of Transactions, Total Sum)

### 2. Query 2

To execute the file for the second query:

1. Start Hadoop and Pig: > start-all.sh
2. Clear output folder for query's result:  
> hadoop fs -rmr  
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-temp  
> hadoop fs -rmr /home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-final
3. Be sure to change the path of files (customer and transaction dataset) in the query2.pig file.
4. Run script in batch mode: > pig PATH/query2.pig
5. The final result will be in folder

/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-final

The script file for this query is below:

```
%declare dataSource1
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Customer'
%declare dataSource2
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Transaction'
%declare tempOutputFile
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-temp'
%declare finalOutputFile
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-final'

raw1 = LOAD '$dataSource1' USING PigStorage(',') AS (CID1, CName,
Age, CountryCode, Salary);

raw2 = LOAD '$dataSource2' USING PigStorage(',') AS (TID, CID2,
TransTotal, TransNumItem:INT, TransDesc);

sub_raw1 = FOREACH raw1 GENERATE CID1, CName, Salary;

joinedTable = JOIN sub_raw1 BY CID1, raw2 BY CID2;

groupedTable = GROUP joinedTable BY CID1;

temp = FOREACH groupedTable GENERATE group, COUNT(joinedTable.CID1)
AS NumOfTrans, SUM(joinedTable.TransTotal) AS TotalSum,
MIN(joinedTable.TransNumItem) AS MinItems;

outputTable = JOIN temp BY group, sub_raw1 BY CID1;

STORE outputTable INTO '$tempOutputFile' USING PigStorage(',');

raw3 = LOAD '$tempOutputFile' USING PigStorage(',') AS (CID,
NumOfTrans, TotalSum, MinItems, CID1, CName, Salary);

outputTable2 = FOREACH raw3 GENERATE CID, CName, Salary, NumOfTrans,
TotalSum, MinItems;
```

```
STORE outputTable2 INTO '$finalOutputFile' USING PigStorage(',');
```

The query's result is attached as file 'query2result.txt'.

The system output (statistics part) is attached as file 'query2SystemOutput.txt'.

Sample output:

```
1,ytchphbgotjz,3919.5823,105,52212.239979000005,1
2,pmjmvymuxbihpywmarw,2588.393,122,60081.70648599999,1
3,wmcfxpmmpp,2402.273,88,45740.707844000004,1
4,slhhzzgyjuw,7308.4614,115,58245.113919999996,1
5,mdudkxpehkmxmuyvvj,5025.3486,116,61284.218937,1
6,wmesdgneqqvnaauv,1173.6986,103,47252.700724999995,1
7,gausenitzw,7229.5146,112,58276.20498099999,1
8,nfnafodqts,6651.0283,93,46487.934261999995,1
9,opcqusrwigbfqomrzy,2851.3687,111,53785.31017199998,1
10,laaxvibivfhfyweh,9814.074,91,44401.128201000014,1
```

Description: (Customer ID, CustomerName, Salary, Number of Transactions, Total Sum, MinItems)

### 3. Query 3

To do query 3, here is the process:

1. We should use two data set (customer and transaction), so we should first join those two data into one variable based on Customer ID.
2. After we have joined data, we group them based on Customer ID and Country Code and aggregate transaction amount using SUM, resulting output of [Customer ID, Country Code, Total Transaction].
3. We perform another separated query of grouping customer data by Country Code and outputting them in format of [Country Code, Total Customer].
4. We join query result from no 3 and no 2 by Country Code. We then group them by Country Code and Total Customer, so we can get the aggregate function of MAX and MIN of the total transaction amount and resulting the final output of [Country Code, Total Customer, Maximum Total Transaction, Minimum Total Transaction].
5. Store the result into file in HDFS.

Query:

```
SET DEFAULT PARALLEL 500;
trans = LOAD '/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/transactioninput'
USING PigStorage(',') as (transid:int,custid:int,
trxamt:float,noitem:int,desc:chararray);
cust = LOAD '/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/customerinput'
USING PigStorage(',') as (id:int,name:chararray, age:int,cc:int,salary:float);
A = join cust by id, trans by custid;
AA = group A by ($0,$3);
AAA = foreach AA generate FLATTEN(group) as (custid, cc), SUM(A.trxamt) as trx;
B = group cust by cc;
```

```

BB = foreach B generate group as cc, COUNT(cust.id) as totalcust;
AB = join AAA by cc, BB by cc;
ABX = group AB by ($1,$4);
C = foreach ABX generate FLATTEN(group) as (cc,totalcust), MIN(AB.trx) ,
MAX(AB.trx);
STORE C INTO
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/project2/query3outputs' USING
PigStorage();

```

Screen captures:

```

grunt> SET DEFAULT_PARALLEL 500;
grunt> trans = LOAD '/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/transaction
input' USING PigStorage(',') as (transid:int,custid:int, trxamt:float,noitem:int
,desc:chararray);
grunt> cust = LOAD '/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/customerinput'
USING PigStorage(',') as (id:int,name:chararray, age:int,cc:int,salary:float)
;
grunt> A = join cust by id, trans by custid;
grunt> AA = group A by ($0,$3);
grunt> AAA = foreach AA generate FLATTEN(group) as (custid, cc), SUM(A.trxamt) a
s trx;
grunt> B = group cust by cc;
grunt> BB = foreach B generate group as cc, COUNT(cust.id) as totalcust;
grunt> AB = join AAA by cc, BB by cc;
grunt> ABX = group AB by ($1,$4);
grunt> C = foreach ABX generate FLATTEN(group) as (cc,totalcust), MAX(AB.trx), M
IN(AB.trx);

```

```

2015-02-22 18:10:52,047 [main] INFO org.apache.pig.tools.pigstats.ScriptState
- Pig features used in the script: HASH_JOIN,GROUP_BY
2015-02-22 18:10:53,011 [main] INFO org.apache.pig.backend.hadoop.executioneng
ine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic?
false
2015-02-22 18:10:53,189 [main] INFO org.apache.pig.backend.hadoop.executioneng
ine.mapReduceLayer.CombinerOptimizer - Choosing to move algebraic foreach to
combiner
2015-02-22 18:10:53,215 [main] INFO org.apache.pig.backend.hadoop.executioneng
ine.mapReduceLayer.CombinerOptimizer - Choosing to move algebraic foreach to
combiner

```

....

```

Input(s):
Successfully read 50000 records (1792839 bytes) from: "/home/ubuntu/Workspace/
hadoop-1.1.0/hadoop-data/customerinput"
Successfully read 5000000 records from: "/home/ubuntu/Workspace/hadoop-1.1.0/h
adoop-data/transactioninput"

Output(s):
Successfully stored 10 records (431 bytes) in: "/home/ubuntu/Workspace/hadoop-
1.1.0/hadoop-data/project2/query3outputs"

Counters:
Total records written : 10
Total bytes written : 431
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_201502181312_0042  ->      job_201502181312_0043,
job_201502181312_0043  ->      job_201502181312_0044,
job_201502181312_0044  ->      job_201502181312_0045,
job_201502181312_0045  ->      job_201502181312_0046,
job_201502181312_0046

2015-02-22 19:12:16,046 [main] INFO  org.apache.pig.backend.hadoop.executionen
gine.mapReduceLayer.MapReduceLauncher - Success!
grunt>

```

Output:

1	5017	31074.71094894409	71878.46228408813
2	4952	33037.3843421936	70566.1169872284
3	4988	32585.36799812317	75037.75818061829
4	4969	28243.546305656433	71899.7768650055
5	4942	31215.28261947632	72883.00854873657
6	5088	32576.66389465332	76930.16226959229
7	5043	31528.636875152588	71983.4383277893
8	5073	31442.272380828857	71460.10570907593
9	5046	32296.86143875122	72303.17418956757
10	4882	29607.029012680054	72920.12038040161

Description: Country Code, Number of Customers, Minimum Total Transaction, Maximum Total Transaction

#### 4. Query 4

The process of execution of this query is the same as query2.

The script for this query is below:

```

%declare dataSource1
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Customer'
%declare dataSource2
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Transaction'

```

```

%declare tempOutputFile
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query4-temp'
%declare finalOutputFile
'/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query4-final'

raw1 = LOAD '$dataSource1' USING PigStorage(',') AS (CID1, CName,
Age, CountryCode, Salary);

raw2 = LOAD '$dataSource2' USING PigStorage(',') AS (TID, CID2,
TransTotal, TransNumItem:INT, TransDesc);

sub_raw1 = FOREACH raw1 GENERATE CID1, CName, Salary;

joinedTable = JOIN raw2 BY CID2, sub_raw1 BY CID1 USING 'replicated';

groupedTable = GROUP joinedTable BY CID1;

temp = FOREACH groupedTable GENERATE group, COUNT(joinedTable.CID1)
AS NumOfTrans, SUM(joinedTable.TransTotal) AS TotalSum,
MIN(joinedTable.TransNumItem) AS MinItems;

outputTable = JOIN temp BY group, sub_raw1 BY CID1 USING
'replicated';

STORE outputTable INTO '$tempOutputFile' USING PigStorage(',');

raw3 = LOAD '$tempOutputFile' USING PigStorage(',') AS (CID,
NumOfTrans, TotalSum, MinItems, CID1, CName, Salary);

outputTable2 = FOREACH raw3 GENERATE CID, CName, Salary, NumOfTrans,
TotalSum, MinItems;

STORE outputTable2 INTO '$finalOutputFile' USING PigStorage(',');

```

The query's result is attached as file 'query4result.txt'.

The system output (statistics part) is attached as file 'query4SystemOutput.txt'.

The statistics from query2 is:

HadoopVersion	PigVersion	UserId	StartedAt	FinishedAt	Features
1.1.0.10.0	ubuntu	2015-02-22	04:55:24	2015-02-22	04:58:35

HASH\_JOIN, GROUP\_BY

Success!

Job Stats (time in seconds):

JobId	Maps	Reduces	MaxMapTime	MinMapTime	AvgMapTime	MaxReduceTime	MinReduceTime	AvgReduceTime	Alias	Feature	Outputs
job_201502220216_0030	1	0	2	2	2	0	0				
0raw1,sub_raw1 MAP_ONLY											
job_201502220216_0031	6	1	10	0	7	57	57	57			
joinedTable,raw2 HASH_JOIN											
job_201502220216_0032	8	1	7	3	5	41	41	41			
groupedTable,temp GROUP_BY,COMBINER											
job_201502220216_0033	2	1	1	1	1	10	10	10			
outputTable HASH_JOIN											
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-temp,											
job_201502220216_0034	1	0	2	2	2	0	0				
0outputTable2,raw3 MAP_ONLY											
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query2-final,											

The statistics from query4 is:

HadoopVersion	PigVersion	UserId	StartedAt	FinishedAt	Features
1.1.0 0.10.0	ubuntu	2015-02-22	05:33:57	2015-02-22	05:36:29

REPLICATED\_JOIN, GROUP\_BY

Success!

Job Stats (time in seconds):

JobId	Maps	Reduces	MaxMapTime	MinMapTime	AvgMapTime	MaxReduceTime	MinReduceTime	AvgReduceTime	Alias	Feature	Outputs
job_201502220216_0053	1	0	2	2	2	0	0				
0raw1,sub_raw1 MULTI_QUERY,MAP_ONLY											
job_201502220216_0054	5	1	26	12	20	86	86	86			
groupedTable,joinedTable,outputTable,raw2,temp											
REPLICATED_JOIN, GROUP_BY, COMBINER											
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query4-temp,											
job_201502220216_0055	1	0	2	2	2	0	0				
0outputTable2,raw3 MAP_ONLY											
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Query4-final,											

Compared the system output of this query and query2, we can clearly see query2 consists of 5 Map and/or MapReduce jobs, whereas query4 requires 3 Map and/or MapReduce jobs. Also, the time cost for query2 & 4 is 3mins 11secs and 2mins 42secs respectively. Using replicated JOIN increase the efficiency by nearly 15%.

Sample output:

```

1,ytchphbgotjz,3919.5823,105,52212.239979,1
2,pmjmvymuxbihpywmrw,2588.393,122,60081.706486,1
3,wmcfxpmmpp,2402.273,88,45740.707844000004,1
4,slhhzzgyjuw,7308.4614,115,58245.113919999996,1
5,mdudkxpehkmxmuyvvj,5025.3486,116,61284.218937000005,1
6,wmesdgneqqvnaauv,1173.6986,103,47252.700725,1
7,gausenitzw,7229.5146,112,58276.204981,1
8,nfnafodqts,6651.0283,93,46487.93426200001,1
9,opcqusrwigbfqomrzy,2851.3687,111,53785.310172000005,1
10,laaxvibivfhfyweh,9814.074,91,44401.128201,1

```

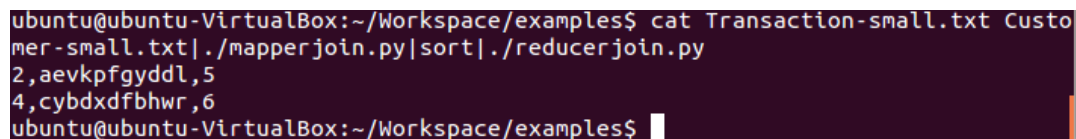
Description: (Customer ID, CustomerName, Salary, Number of Transactions, Total Sum, MinItems)

## 5. Query 5

We choose Python language combined with the Hadoop Streaming. To do the query, we need one mapper (mapperjoin.py) and one reducer (reducerjoin.py), where the joining data set will happen in reducer side. The mapper will read the input from both of transaction and customer file. Initial value of (-1) is set for each of the attribute that we need (customer\_id, customer\_name, country\_code, transaction) and will be replaced with real value after reading the input file. So for example, when mapper read data from customer file, the value of transaction\_id will still be -1 while other attributes will have the real values. The output from the mapper will follow the format of [customer\_id, customer\_name, country\_code, transaction\_id]. The output will be sorted and then sent to reducer. Reducer will group the output from the mapper based on **customer ID**, filter the result using condition (if country code = 5) and then sent to stdout.

We first test the codes using the local Python environment first using a small dataset:

```
cat Transaction-small.txt Customer-small.txt|./mapperjoin.py|sort|./reducerjoin.py
```



```

ubuntu@ubuntu-VirtualBox:~/Workspace/examples$ cat Transaction-small.txt Customer-small.txt|./mapperjoin.py|sort|./reducerjoin.py
2,aevkpfgyddl,5
4,cybdxdfbhwr,6
ubuntu@ubuntu-VirtualBox:~/Workspace/examples$

```

After we make sure it works, then we run the mapper and reducer script using Hadoop with the real data set using following command:

```

$HADOOP_HOME/bin/hadoop jar /home/ubuntu/Workspace/hadoop-1.1.0/contrib/streaming/hadoop-streaming-1.1.0.jar -D mapred.reduce.tasks=1 -file /home/ubuntu/Workspace/examples/python/mapperjoin.py -mapper /home/ubuntu/Workspace/examples/python/mapperjoin.py -file /home/ubuntu/Workspace/examples/python/reducerjoin.py -reducer /home/ubuntu/Workspace/examples/python/reducerjoin.py -input /user/ubuntu/inputbig/* -output /user/ubuntu/outputbig

```

Screen capture:



```

15/02/22 00:53:29 INFO streaming.StreamJob: map 100% reduce 91%
15/02/22 00:53:31 INFO streaming.StreamJob: map 100% reduce 92%
15/02/22 00:53:32 INFO streaming.StreamJob: map 100% reduce 93%
15/02/22 00:53:33 INFO streaming.StreamJob: map 100% reduce 94%
15/02/22 00:53:35 INFO streaming.StreamJob: map 100% reduce 95%
15/02/22 00:53:36 INFO streaming.StreamJob: map 100% reduce 96%
15/02/22 00:53:38 INFO streaming.StreamJob: map 100% reduce 97%
15/02/22 00:53:39 INFO streaming.StreamJob: map 100% reduce 98%
15/02/22 00:53:40 INFO streaming.StreamJob: map 100% reduce 99%
15/02/22 00:53:43 INFO streaming.StreamJob: map 100% reduce 100%
15/02/22 00:53:45 INFO streaming.StreamJob: Job complete: job_201502181312_0039
15/02/22 00:53:45 INFO streaming.StreamJob: Output: /user/ubuntu/outputbig

```

Notice that we set the number of reducer to be equal 1 in the above command (-D mapred.reduce.tasks=1). This is because by default, hadoop streaming will use 3 reducers which will divide the result and group the count result separately in each file. So we need to specify the number of reducer to be in a single reducer so that the result will be grouped in one single output.

The query resulted **4942 records** (customer which has country code = 5).

Sample output:

```

100,mrslucbrmyyibz,121
10005,hgaxmvsgccudwnjf,116
10012,qarmxordkaznloacf,119
10018,zinwkicxrmdgz,94
10019,xcwgnorgbamsftl,89
10029,izrgghtbeewv,119
1007,yumqzkddbyrgwexc,93
10071,eimxiflvito,105
10077,ibzbyctabz,99
10079,ewnhyjnbvg,85
10090,iqzxcrlnnqdcrdh,112
10100,oevswbfbmruzwmrm,108
10117,erllnxedhwo,121

```

Description: Customer ID, Customer Name, Count Transactions

## 6. Query 6

Still, we choose Python to do this query. The mapper function is attached as file 'PMapper.py' and the reducer function is 'PReducer.py'.

To execute the file for the sixth query:

1. Start Hadoop and Pig
2. Clear output folder for query's result: > `hadoop fs -rmr /home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/query6`
3. Go to hadoop local directory.
4. Make sure our Mapper and Reducer function is executable.
5. Specify the number of reducer to be 1 and run: > `bin/hadoop jar contrib/streaming/hadoop-streaming-1.1.0.jar -D mapred.reduce.tasks=1 -file /home/ubuntu/Desktop/PMapper.py -mapper /home/ubuntu/Desktop/PMapper.py -file /home/ubuntu/Desktop/PReducer.py -reducer /home/ubuntu/Desktop/PReducer.py -input`

```
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/Transaction -output  
/home/ubuntu/Workspace/hadoop-1.1.0/hadoop-data/query6
```

The reason of specifying the reducer number to be 1 is that, the reducers input under the Hadoop streaming mode is just a stream. These means all the Key-Value pairs will fit into all reducers, instead of each Key-Value pair goes to different reducer. So, we limit the number of reducer in order to fit all our result in one file, which is attached as 'query6result.txt'.

Compared to query1, query6 is more Hadoop visible, more close to MapReduce algorithm, and query1 is more like a SQL language. Also, as described above the input of reducer in query6 is totally different from typical JAVA MapReduce framework.

In a word, the Hadoop streaming mode makes our MapReduce algorithm even difficult to perform the SQL queries. Running in streaming mode, it takes 48 seconds to finish this query, whereas in query1, it takes 58 seconds. So, streaming in Python is no more convenient than JAVA but is more efficient than Pig.

Sample output:

```
1,105,52212.239979  
10,91,44401.128201  
100,115,57376.429365  
1000,105,58084.060554  
10000,87,42559.973868  
10001,112,54977.4128525
```

Description: CustomerID, NumTransactions, TotalSum